

Personas

Sarah Mitchell

Age: 28

Occupation: Marketing Manager

Background: Extremely busy professional with long, irregular work hours. Wants to get in shape for her wedding but struggles to find time for gym research and meal prep. Needs an automated, efficient solution to remove guesswork and save time.

David Rodriguez

Age: 42

Occupation: High School Teacher

Background: Former athlete who has become sedentary and gained weight. Overwhelmed by conflicting fitness information online and doesn't know where to start. Needs a structured, guided plan to rebuild healthy habits and lose 30 pounds.

James O'Reilly

Age: 55

Occupation: Accountant

Background: Recently had a health scare and was advised by his doctor to improve his diet and exercise. Concerned about performing exercises with correct form to avoid injury. Needs a safe, personalized plan that accounts for his age and specific health parameters.

Chloe Bennett

Age: 23

Occupation: Fitness Influencer

Background: Very knowledgeable about fitness but spends excessive time creating her own content and plans. Needs to automate her routine planning to free up time for content creation. Wants to easily share her progress and data with her engaged online community.

Emily Carter

Age: 31

Occupation: Software Developer

Background: Prefers data-driven approaches and loves tracking metrics. Lacks motivation to exercise consistently without external accountability. Wants a highly personalized plan based on her biometrics and needs to share her progress with friends for motivation and friendly competition.

Discussion Output

David: Okay, team, let's break this down. So, we're looking at an app that supposedly creates a custom workout and diet plan based on my body stats and goals. My first question: what *specific* biometric data are we talking? Just height, weight, and age? Or does it connect to my smartwatch for heart rate, my scale for body fat percentage...? The quality of the plan is entirely dependent on the quality of the input data.

Chloe: That's a good point, David. For me, simplicity is key. I don't have time to manually input twenty data points every morning. If it can sync with my Apple Health app and pull my step count, sleep data, and weight from my smart scale automatically, that would be a game-changer. But more importantly, when it says "share progress with friends," does that mean it creates a leaderboard? Because as a mom, I don't need more competition; I need a support group to cheer me on when I just manage to drink enough water for the day.

Robert: I have to agree with Chloe on the simplicity front. My doctor told me to move more and watch my salt intake. This app sounds... complex. Will it ask me about my arthritis or my previous knee surgery? I'm not interested in bodybuilding; I need a safe walking schedule and maybe some simple exercises I can do in my living room. How will it know what "easy-to-follow" means for someone my age?

Sophia: See, for me, "easy-to-follow" is useless. I need advanced periodization, macronutrient cycling, and the ability to adjust my caloric intake based on my daily energy expenditure, which changes with my teaching schedule. Robert, to your point, the app *must* have a comprehensive onboarding questionnaire that covers injuries, mobility limitations, and experience level. David, it should absolutely take body fat percentage; it's a useless plan for body composition without it. And Chloe, the social sharing for someone like me is vital—it's my workout log and a platform for accountability with my competition peers.

Ben: As a physical therapist, Sophia's last point is the most critical to me. Form. How does the app ensure I'm doing the exercises correctly to protect my shoulder? Will it have video demonstrations? AI form-checking through the phone's camera? If it just gives me a list of exercises and I perform them wrong, it's worse than doing nothing. Furthermore, when I input my injury, will it intelligently substitute exercises? For example, if the "ideal" shoulder exercise is an overhead press, but my injury can't

handle that, will it automatically suggest a landmine press or a resistance band alternative?

David: Ben, that's brilliant. I didn't even think about exercise substitution. That's a huge algorithmic challenge. So, we're establishing that the app needs a deep, multi-layered input: biometrics (basic and advanced, with sync capabilities), goals, injury history, mobility constraints, *and* daily schedule availability. My next question: how dynamic is the plan? If I have a hell week at work and only sleep 5 hours a night, will it adjust my expected recovery and maybe deload the workouts automatically? Or is it a static plan?

Chloe: Yes! And on the nutrition side, if I say "sustainable weight loss," does it give me a generic 1200-calorie plan? Or does it consider that I have to cook for two kids and might need "quick, healthy family meal" suggestions? I need it to work with my life, not against it. The sharing feature would be perfect if I could just share my weekly "win" (like "I hit my step goal 5/7 days!") with my sister for a virtual high-five, not my daily calorie count.

Robert: This is becoming clearer. So, it's not just a plan generator; it's an adaptive coach. It should ask me how I'm feeling each day. If I report fatigue, it might suggest a shorter walk. And the sharing for me would be with my daughter, so she can see I'm being active and not worry. But I'm concerned about cost. An app this sophisticated, with AI and video demonstrations... surely it's a pricey monthly subscription?

Sophia: It has to be tiered. A free version for basic goals, a premium for advanced analytics and custom exercise libraries, and maybe a coaching tier where you can get actual human oversight. For my bodybuilding prep, I'd pay for that. But back to the nutrition: David mentioned it was a "nutritional plan." Does that mean it generates a meal plan? Or just gives macro targets? I need the former, with the ability to swap out foods I dislike.

Ben: And for me, the nutritional plan must respect medical conditions. If Robert has pre-hypertension, it shouldn't be suggesting high-sodium foods. The app needs to be more than a fitness tool; it needs a layer of medical awareness, even if it's just a disclaimer to consult your doctor.

David: So, to synthesize: We're envisioning an intelligent, adaptive platform. It uses comprehensive user data (manually input and synced) to generate a personalized, safe workout regimen with exercise substitutions and form guidance. It creates a tailored nutritional strategy, from simple calorie goals to full meal plans, considering user constraints. It then adjusts these plans dynamically based on user feedback and progress tracking. Finally, it has a flexible social sharing feature that allows for everything from basic accountability cheerleading to showcasing a detailed bodybuilder's log.

Chloe: That sounds... almost too good to be true. But if it existed, and it was intuitive, it would solve everyone's problems in this room. We all just want a plan that *fits* us, without the overwhelming guesswork.

Robert: Well, when you put it that way, David, I'm less intimidated. It sounds less like a rigid taskmaster and more like a helpful assistant. As long as it remembers I'm 65 and not 25.

Sophia: Exactly. The core of the idea is *customization*. A one-size-fits-all approach fails for everyone. This software's success hinges on the depth of its personalization algorithms.

Ben: And its commitment to safety and proper movement. Without that, it's just a fancy list generator. But with it? It could be genuinely transformative.

Persona Dialogue

David: Alright team, we've agreed the core value is in a smart, adaptive algorithm. But we can't build that in one go. We need to break this down into sprints. I think our first sprint has to be about creating the foundation. We need a secure user profile that captures all the critical initial data: biometrics, goals, available time, equipment, injuries, and dietary restrictions. Without that, we have nothing to base the algorithm on.

Chloe: Absolutely. And it has to be dead simple. A clean, intuitive onboarding process. If I get frustrated filling out a long form, I'm abandoning the app before I even start. This first sprint should also include a very basic, static version of the plan generator. Nothing fancy, just proof that we can take the inputs and produce a simple workout and diet suggestion.

Robert: I agree, but safety first. That basic plan generator must have hard-coded rules to avoid suggesting dangerous exercises for common conditions like knee arthritis or heart issues. We can start with a limited exercise library that only includes safe, beginner-friendly options.

Ben: As a physical therapist, I have to stress that the exercise library is paramount, even in V1. Each exercise needs a clear name, a description of the movement, and most importantly, a video demonstration or at the very least, high-quality GIFs showing proper form. We cannot ship anything without this. It's a non-negotiable safety feature.

Sophia: I can see the logic in starting simple, but we must architect the user profile and plan data structures with extreme customization in mind from the very beginning. The database schema needs to handle complex future needs, like weekly changing

macronutrient targets and detailed training splits, even if we don't build those features now. Otherwise, we'll have to refactor everything later.

David: Good point, Sophia. Tech debt early on would kill us. So Sprint 1: Robust User Onboarding, a foundational (but limited) exercise library with form visuals, and a bare-bones plan generator that uses simple, safe heuristics.

Chloe: Now, for Sprint 2, I think we need the tracking and feedback loop. What's the point of a plan if you can't log that you've done it? This sprint needs to let users check off completed workouts, log their weight, and maybe track basic nutrition. This is how we start gathering the data for the "adaptive" part.

Robert: And a simple notification system for those gentle reminders. "Time for your walk, Robert." This is crucial for users like me who are building a new habit.

Ben: This is also where we introduce the first layer of adaption. If a user consistently logs "could not complete" or "felt pain during" for a specific exercise, the app should be smart enough to suggest a predefined regression or alternative from the library. It's not full AI learning yet, but it's a crucial step.

Sophia: I'd also push for basic sharing in this sprint. Even if it's just the ability to send a simple "I completed my week!" message to a designated contact. This addresses the accountability need for Chloe and Robert immediately, before we build complex social feeds.

David: Okay, Sprint 2: Activity Logging, Basic Notifications, Simple Adaptive Feedback (exercise substitutions), and Core Sharing Functionality.

Sophia: Now for Sprint 3, we can tackle the advanced customization I need. This is where we build out the interfaces for users to fully customize their training splits (e.g., PPL, Upper/Lower), set specific macronutrient goals, and create more complex dietary preferences like vegan or keto. The algorithm now has to be smart enough to generate plans based on these highly specific parameters.

Ben: This sprint should also expand our exercise library significantly and introduce the ability for users to build a "restriction list" – to manually flag exercises they cannot do due to injury, which the algorithm will then permanently avoid.

Chloe: As long as the default remains simple for users like me who don't want to fiddle with these settings. But I agree, the power needs to be there for those who want it.

David: And this is probably where we need to start integrating with third-party apps like MyFitnessPal for easier nutrition tracking. Sprint 3: Advanced Customization Tools, Expanded Exercise & Restriction Management, and Initial Third-Party Integrations.

Robert: What comes next? The app is functional, but how does it become that "daily partner" Ben described?

Ben: Sprint 4 is for the true AI learning. We use all the data we've been collecting—completed sets, logged weights, user feedback on difficulty and pain—to make the algorithm predictive. It should now proactively suggest increasing weights, deload weeks, or entirely new exercise variations to break plateaus. This is the "revolutionary" feedback loop.

Sophia: And this is where we can flesh out the social features into a true community platform. Detailed progress sharing, the ability to create groups, and share whole workout plans. Granular privacy controls are a must here.

Chloe: For me, the killer feature in a later sprint would be calendar sync and automated schedule adjustment. But that feels like a polish item once the core AI is solid.

David: Right. So let's call Sprint 4: AI-Powered Adaptation & Advanced Social Community. Then we can have a final Sprint 5 for polish, performance, and those deeper integrations like calendar sync and maybe even wearable device data.

Robert: That sequence makes sense. We build the foundation, then teach it to learn from us, and finally make it seamless to use in our daily lives.

Sprints

Sprint 1: Foundation & Safety Onboarding

Duration: 3 weeks

Card seleccionado: #7 Privacy and Data

Goals: Create user profile system, Build basic exercise library with form visuals, Develop static plan generator with safety rules

Tasks: Design user onboarding flow, Implement biometric and goal data capture, Develop limited exercise database with video/GIF demonstrations, Build heuristic-based plan generator with age/condition safeguards

Sprint 2: Core Tracking & Feedback

Duration: 2 weeks

Card seleccionado: #10 Human Agency

Goals: Implement activity logging, Enable basic adaptive feedback, Introduce core sharing and notification features

Tasks: Develop workout and nutrition logging modules, Build notification/reminder system, Create simple exercise substitution logic, Implement basic progress sharing to contacts

Sprint 3: Advanced Customization

Duration: 3 weeks

Card seleccionado: #15 Stakeholder Participation

Goals: Deliver tools for advanced users, Expand exercise and diet options, Begin third-party integrations

Tasks: Build custom training split editor, Develop detailed macronutrient and dietary preference settings, Significantly expand exercise library, Implement user injury restriction lists, Develop API for MyFitnessPal integration

Sprint 4: AI Learning & Community

Duration: 3 weeks

Card seleccionado: #7 Privacy and Data

Goals: Implement predictive algorithm adjustments, Build advanced social sharing platform

Tasks: Develop AI engine for progressive overload and deload recommendations, Create community feed and group features, Implement granular privacy controls for sharing, Enhance algorithm based on aggregated user feedback data

Sprint 5: Polish & Deep Integration

Duration: 2 weeks

Card seleccionado: #0 Stakeholder Analysis

Goals: Refine user experience, Integrate with external calendars and devices, Perform performance optimization

Tasks: Implement calendar sync for automatic schedule adjustment, Develop integration with health platforms (Apple Health, Google Fit), Conduct performance profiling and optimization, Polish UI/UX based on user testing feedback

ECCOLA Cards Selected per Sprint

Comparação entre a seleção dos cards:

Sprint 1- Privacidade e dados em comum

Sprint 2- Human agency em comum

Sprint 3- Não teve em comum

Sprint 4- Privacidade e data em comum

Sprint 5- Human agency em comum



Sprint 1: User Foundation & Exercise Library

Duration: 3 weeks

Goals: Create a comprehensive user onboarding process, Build a secure user profile database, Develop a searchable exercise library with video demonstrations and constraint tags

Tasks: Design and implement user onboarding questionnaire, Develop database schema for user profiles (biometrics, goals, injuries, restrictions), Create exercise database with video content, Tag exercises with modifiers (e.g., low-impact, knee-friendly, shoulder-safe)



Stakeholder Analysis (#0 - Analyze)

Motivation: In order to understand the big picture, it is important to first understand who the system can affect and how. Try to also think past the obvious, direct stakeholders such as your end-users.

What to Do:

- Identify stakeholders.
- Who does the system affect and how? Stakeholders are not simply users, developers and customers.
- How are the various stakeholders linked together?
- Can these different stakeholders influence the development of the system? How?
- Remember that a user is often an organization and the end-user is an individual. Similarly, AI systems can treat people as objects for data collection.

Practical Example: Autonomous cars don't just affect their passengers. Anyone nearby is affected; some even change the way they drive. If at one point half of the traffic consists of self-driving cars, what are the societal impacts of such systems? E.g., how are the people who can't afford one affected? Regulations arising from such systems also affect everyone.

Justification: This sprint involves collecting user data (biometrics, goals, injuries) and building an exercise library, which affects not only end-users but also healthcare providers, family members (due to sharing features), and regulatory bodies. Understanding all stakeholders helps ensure inclusive and ethical development.

Communication (#3 - Transparency)

Motivation: In practice, communication is a big part of being transparent with your stakeholders. Being transparent in communication can generate trust.

What to Do:

- Ask yourself:
- What is the goal of the system? Why is this particular system deployed in this specific area?
- What do you communicate about the system to its users and end-users? Is it enough for them to understand how the system works?
- If relevant to your system, do you somehow tell your (end-)users that they are interacting with an AI system and not with another human being?
- Do you collect user feedback? How is it used to change/improve the system?
- Are communication and transparency towards other audiences, such as the general public, relevant?

Practical Example: Clearly stating what data you collect and why can make you seem much more trustworthy. Compare this to a cellphone application that just states it needs to access your camera and storage.

Justification: During user onboarding, clear communication about data collection (biometrics, injuries), how it's used for personalized plans, and sharing features is crucial for building trust and obtaining informed consent.

Privacy and Data (#7 - Data)

Motivation: Privacy is a rising trend in the wake of various recent data misuse reveals. People are now increasingly conscious about handing out personal data. Similarly, regulations such as GDPR now affect data collection.

What to Do:

- Ask yourself:
- What data are used by the system?

- Does the system use or collect personal data? Why? How is the personal data used?
- Do you clearly inform your (end-)users about any personal data collection? E.g., ask for consent, provide an opportunity to revoke it etc.
- Have you taken measures to enhance (end-user) privacy, such as encryption or anonymization?
- Who makes the decisions regarding data use and collection? Do you have organizational policies for it?

Practical Example: Rather than collecting and selling data, appealing to privacy can also be profitable. Regulations are making it increasingly difficult to collect lots of personal data for profit. Privacy can be an alternate selling point in today's climate.

Justification: The sprint involves collecting sensitive biometric and health data (injuries, restrictions) in user profiles, raising privacy concerns that must be addressed through secure storage, clear consent mechanisms, and compliance with data protection laws.

Data Quality (#8 - Data)

Motivation: As AI are trained using data, the data used directly affects how the system operates. Both the nature and the quality and integrity of the data used has to align with goals of the system.

What to Do:

- Ask yourself:
- What are good or poor quality data in the context of your system?
- How do you evaluate the quality and integrity of your own data? Are there alternative ways?
- If you utilize data from external sources, how do you control their quality?
- Did you align your system with relevant standards (for example ISO, IEEE) or widely adopted protocols for daily data management and governance?
- How can you tell if your data sets have been hacked or otherwise compromised?

Practical Example: In 2017, Amazon scrapped its recruitment AI because of bad data. They used past recruitment data to teach the AI. As they had mostly hired men, the AI began to consider women undesirable based on the data.

Justification: The exercise library with video content and constraint tags must be accurate and reliable to prevent harm (e.g., recommending unsafe exercises for users with injuries), making data quality critical for user safety and trust.

Accessibility (#14 - Fairness)

Motivation: Technology can be discriminatin in various ways. Given the enormous impact AI systems can have, ensuring equal access to their positive impacts is ethically important.

What to Do:

- Ask yourself:
- Does the system consider a wide range of individual preferences and abilities? If not, why?
- Is the system usable by those with special needs or disabilities, those at risk of exclusion, or those using assistive technologies?
- Were people representing various groups somehow involved in the development of the system?
- How is the potential user audience taken into account?
- Is the team involved in building the sustem representative of your largel user audience? Is it representative of the general population?
- Did you assess whether there could be (groups of) people?

Practical Example: AI tends to benefit those who are already technologically capable, resulting in increased inequality. E.g. most of the images used in machine learning have been labeled by young white men.


Justification: The sprint includes tagging exercises with modifiers (e.g., low-impact, knee-friendly) to cater to users with different abilities and restrictions, highlighting the need for accessibility to ensure the system is inclusive and fair.

Sprint 2: Core Plan Generation

Duration: 2 weeks

Goals: Develop the initial plan generation algorithm, Deliver basic workout and nutrition plans to the user

Tasks: Build algorithm to generate workout plans based on user profile, Build algorithm to generate basic nutrition guidelines based on profile, Display generated plans in the user interface, Implement basic print/save functionality

-  Stakeholder Analysis (#0 - Analyze)
-  Explainability (#2 - Transparency)
-  Privacy and Data (#7 - Data)
-  Data Quality (#8 - Data)
-  Human Agency (#10 - Agency & Oversight)

Sprint 3: Feedback & Adaptation Loop

Duration: 3 weeks

Goals: Enable users to log completed activities and provide feedback, Implement basic plan adaptation based on user input, Introduce reminder system for habit formation

Tasks: Develop workout and nutrition logging features, Create feedback mechanism (e.g., rate difficulty, log pain), Build logic for exercise substitution based on feedback, Implement simple notification system for reminders






-  Privacy and Data (#7 - Data)
-  Data Quality (#8 - Data)
-  Explainability (#2 - Transparency)
-  Communication (#3 - Transparency)
-  System Safety (#13 - Safety & Security)

Sprint 4: Social Accountability & Sharing

Duration: 2 weeks

Goals: Implement granular privacy-controlled sharing features, Develop a system for social support and accountability

Tasks: Build friend/connection system, Create shareable progress dashboards with customizable data visibility, Develop notification system for accountability buddies, Implement community features like comment threads on shared progress

-  Privacy and Data (#7 - Data)
-  Access to Data (#9 - Data)
-  Communication (#3 - Transparency)
-  Human Agency (#10 - Agency & Oversight)
-  Auditability (#18 - Accountability)

Sprint 5: Advanced Customization & Integration

Duration: 3 weeks

Goals: Deliver advanced features for expert users, Enable deeper integration with external tools and data

Tasks: Build advanced plan customizer (e.g., specific training splits, macro targets), Develop API integration with nutrition apps like MyFitnessPal, Enhance adaptation algorithm with more complex AI patterns, Implement data export functionality

-  Explainability (#2 - Transparency)
-  Privacy and Data (#7 - Data)
-  Data Quality (#8 - Data)
-  Access to Data (#9 - Data)
-  Human Agency (#10 - Agency & Oversight)

Are the selected cards aligned with your project goals?



Sprint 2: Core Plan Generation

Duration: 2 weeks

Goals: Develop the initial plan generation algorithm, Deliver basic workout and nutrition plans to the user

Tasks: Build algorithm to generate workout plans based on user profile, Build algorithm to generate basic nutrition guidelines based on profile, Display generated plans in the user interface, Implement basic print/save functionality



Stakeholder Analysis (#0 - Analyze)

Motivation: In order to understand the big picture, it is important to first understand who the system can affect and how. Try to also think past the obvious, direct stakeholders such as your end-users.

What to Do:

- Identify stakeholders.
- Who does the system affect and how? Stakeholders are not simply users, developers and customers.
- How are the various stakeholders linked together?
- Can these different stakeholders influence the development of the system? How?
- Remember that a user is often an organization and the end-user is an individual. Similarly, AI systems can treat people as objects for data collection.

Practical Example: Autonomous cars don't just affect their passengers. Anyone nearby is affected; some even change the way they drive. If at one point half of the traffic consists of self-driving cars, what are the societal impacts of such systems? E.g., how are the people who can't afford one affected? Regulations arising from such systems also affect everyone.

Justification: This sprint involves generating personalized health plans, which directly affects end-users' wellbeing. Stakeholders include users, their friends/family (via sharing features), healthcare providers, nutritionists, fitness experts, and regulatory bodies. Understanding these relationships is crucial for ethical development.



Explainability (#2 - Transparency)

Motivation: If we cannot understand the reasons behind the actions of the AI, it is difficult to trust it.

What to Do:

- Ask yourself:

- Is explainability a goal for your system? How do you plan to ensure it?
- How well can each decision of the system be understood? By both developers and (end-)users?
- Did you try to use the simplest and most interpretable model possible for the context?
- Did you make trade-offs between explainability and accuracy? What kind of? Why?
- How familiar are you with your training or testing data? Can you change it when needed?
- If you utilize third party components in the system, how well do you understand them?

Practical Example: When interacting with a robot, users could ideally ask the robot 'why did you do that?' and receive an understandable response. This would make it much easier for them to trust a system.

Justification: The algorithm generating workout/nutrition plans must be explainable to users so they understand why specific recommendations are made, especially since these directly impact their health and safety.

Privacy and Data (#7 - Data)

Motivation: Privacy is a rising trend in the wake of various recent data misuse reveals. People are now increasingly conscious about handing out personal data. Similarly, regulations such as GDPR now affect data collection.

What to Do:

- Ask yourself:
- What data are used by the system?
- Does the system use or collect personal data? Why? How is the personal data used?
- Do you clearly inform your (end-)users about any personal data collection? E.g., ask for consent, provide an opportunity to revoke it etc.
- Have you taken measures to enhance (end-user) privacy, such as encryption or anonymization?
- Who makes the decisions regarding data use and collection? Do you have organizational policies for it?

Practical Example: Rather than collecting and selling data, appealing to privacy can also be profitable. Regulations are making it increasingly difficult to collect lots of personal data for profit. Privacy can be an alternate selling point in today's climate.

Justification: This sprint collects sensitive biometric and health data to generate personalized plans. Privacy considerations are paramount, including clear consent mechanisms and data protection measures.

Data Quality (#8 - Data)

Motivation: As AI are trained using data, the data used directly affects how the system operates. Both the nature and the quality and integrity of the data used has to align with goals of the system.

What to Do:

- Ask yourself:
- What are good or poor quality data in the context of your system?
- How do you evaluate the quality and integrity of your own data? Are there alternative ways?
- If you utilize data from external sources, how do you control their quality?
- Did you align your system with relevant standards (for example ISO, IEEE) or widely adopted protocols for daily data management and governance?
- How can you tell if your data sets have been hacked or otherwise compromised?

Practical Example: In 2017, Amazon scrapped its recruitment AI because of bad data. They used past recruitment data to teach the AI. As they had mostly hired men, the AI began to consider women undesirable based on the data.

Justification: The quality of input biometric data and training data for the algorithms directly impacts the safety and effectiveness of generated health plans. Poor data could lead to harmful recommendations.

Human Agency (#10 - Agency & Oversight)

Motivation: People interacting with the system or using it should be able to understand it sufficiently. Users should be able to make informed decisions based on its suggestions, or to challenge its suggestions. AI systems should let humans make independent choices.

What to Do:

- Ask yourself:
- Does the system interact with decisions by human actors, i.e. end users (e.g. recommending users actions or decisions, or presenting options)?
- Does the system communicate to its (end) users that a decision, content or outcome is the result of an algorithmic decision? Into how much detail does it go?
- In the system's use context, what tasks are done by the system and what tasks are done by humans?
- Have you taken measures to prevent overconfidence or overreliance on the system?

Practical Example: A medical system recommends diagnoses. How does the system communicate to doctors why it made a recommendation? How should the doctors know when to challenge the system? Does the system somehow change how patients and doctors interact?

Justification: The system generates health recommendations that users might follow blindly. Ensuring users maintain agency to question, modify, or reject recommendations is critical for safety and ethical use.

Sprint 3: Feedback & Adaptation Loop

Duration: 3 weeks

Goals: Enable users to log completed activities and provide feedback, Implement basic plan adaptation based on user input, Introduce reminder system for habit formation

Tasks: Develop workout and nutrition logging features, Create feedback mechanism (e.g., rate difficulty, log pain), Build logic for exercise substitution based on feedback, Implement simple notification system for reminders






-  Privacy and Data (#7 - Data)
-  Data Quality (#8 - Data)
-  Explainability (#2 - Transparency)
-  Communication (#3 - Transparency)
-  System Safety (#13 - Safety & Security)

Sprint 4: Social Accountability & Sharing

Duration: 2 weeks

Goals: Implement granular privacy-controlled sharing features, Develop a system for social support and accountability

Tasks: Build friend/connection system, Create shareable progress dashboards with customizable data visibility, Develop notification system for accountability buddies, Implement community features like comment threads on shared progress

-  Privacy and Data (#7 - Data)
-  Access to Data (#9 - Data)
-  Communication (#3 - Transparency)
-  Human Agency (#10 - Agency & Oversight)
-  Auditability (#18 - Accountability)

Sprint 5: Advanced Customization & Integration

Duration: 3 weeks

Goals: Deliver advanced features for expert users, Enable deeper integration with external tools and data

Tasks: Build advanced plan customizer (e.g., specific training splits, macro targets), Develop API integration with nutrition apps like MyFitnessPal, Enhance adaptation algorithm with more complex AI patterns, Implement data export functionality

📌 Explainability (#2 - Transparency)

📌 Privacy and Data (#7 - Data)

📌 Data Quality (#8 - Data)

📌 Access to Data (#9 - Data)

📌 Human Agency (#10 - Agency & Oversight)

Are the selected cards aligned with your project goals?

Sprint 3: Feedback & Adaptation Loop

Duration: 3 weeks

Goals: Enable users to log completed activities and provide feedback, Implement basic plan adaptation based on user input, Introduce reminder system for habit formation

Tasks: Develop workout and nutrition logging features, Create feedback mechanism (e.g., rate difficulty, log pain), Build logic for exercise substitution based on feedback, Implement simple notification system for reminders

📌 Privacy and Data (#7 - Data)

Motivation: Privacy is a rising trend in the wake of various recent data misuse reveals. People are now increasingly conscious about handing out personal data. Similarly, regulations such as GDPR now affect data collection.

What to Do:

- Ask yourself:
- What data are used by the system?
- Does the system use or collect personal data? Why? How is the personal data used?
- Do you clearly inform your (end-)users about any personal data collection? E.g., ask for consent, provide an opportunity to revoke it etc.
- Have you taken measures to enhance (end-user) privacy, such as encryption or anonymization?
- Who makes the decisions regarding data use and collection? Do you have organizational policies for it?

Practical Example: Rather than collecting and selling data, appealing to privacy can also be profitable. Regulations are making it increasingly difficult to collect lots of personal data for profit. Privacy can be an alternate selling point in today's climate.

Justification: This sprint involves developing logging features for workout and nutrition, which collect sensitive personal and health data. Ensuring privacy is critical to comply with regulations like GDPR and build user trust.

Data Quality (#8 - Data)

Motivation: As AI are trained using data, the data used directly affects how the system operates. Both the nature and the quality and integrity of the data used has to align with goals of the system.

What to Do:

- Ask yourself:
- What are good or poor quality data in the context of your system?
- How do you evaluate the quality and integrity of your own data? Are there alternative ways?
- If you utilize data from external sources, how do you control their quality?
- Did you align your system with relevant standards (for example ISO, IEEE) or widely adopted protocols for daily data management and governance?
- How can you tell if your data sets have been hacked or otherwise compromised?

Practical Example: In 2017, Amazon scrapped its recruitment AI because of bad data. They used past recruitment data to teach the AI. As they had mostly hired men, the AI began to consider women undesirable based on the data.

Justification: The feedback mechanism (e.g., rating difficulty, logging pain) provides data for plan adaptation. Poor data quality could lead to incorrect adaptations, potentially causing harm or inefficacy in fitness plans.

Explainability (#2 - Transparency)

Motivation: If we cannot understand the reasons behind the actions of the AI, it is difficult to trust it.

What to Do:

- Ask yourself:
- Is explainability a goal for your system? How do you plan to ensure it?
- How well can each decision of the system be understood? By both developers and (end-)users?
- Did you try to use the simplest and most interpretable model possible for the context?
- Did you make trade-offs between explainability and accuracy? What kind of? Why?
- How familiar are you with your training or testing data? Can you change it when needed?

- If you utilize third party components in the system, how well do you understand them?

Practical Example: When interacting with a robot, users could ideally ask the robot 'why did you do that?' and receive an understandable response. This would make it much easier for them to trust a system.

Justification: The logic for exercise substitution based on feedback requires transparency. Users need to understand why changes are made to their plans to trust and effectively use the system.

Communication (#3 - Transparency)

Motivation: In practice, communication is a big part of being transparent with your stakeholders. Being transparent in communication can generate trust.

What to Do:

- Ask yourself:
- What is the goal of the system? Why is this particular system deployed in this specific area?
- What do you communicate about the system to its users and end-users? Is it enough for them to understand how the system works?
- If relevant to your system, do you somehow tell your (end-)users that they are interacting with an AI system and not with another human being?
- Do you collect user feedback? How is it used to change/improve the system?
- Are communication and transparency towards other audiences, such as the general public, relevant?

Practical Example: Clearly stating what data you collect and why can make you seem much more trustworthy. Compare this to a cellphone application that just states it needs to access your camera and storage.

Justification: This sprint involves a feedback mechanism and notifications. Clear communication about data usage, feedback processing, and reminder purposes is essential for user awareness and consent.

System Safety (#13 - Safety & Security)

Motivation: AI systems exert notable influence on the physical world whether they are cyber-physical or not. Various risks and their consequences should be considered, thinking ahead to the operational life of the system.

What to Do:

- Ask yourself:

- What kind of risks does the system involve? What kind of damage could it cause?
- How do you measure and assess risks and safety?
- In what conditions do the fallback plans trigger? Are they automatic or do they require human input?
- Is there a plan to mitigate or manage technological errors, accidents, or malicious misuse? What if the system provides wrong results, becomes unavailable, or provides societally unacceptable results?
- What liability and consumer protection laws apply to your system? Have you taken them into account?

Practical Example: AI systems can aid automating various organizational tasks, making it possible to reduce personnel. However, if a customer organization becomes reliant on your AI system to handle a portion of its operations, what happens if that AI stops functioning for even a few days? What could you do to alleviate the impact?






Justification: Adapting workout plans based on feedback could lead to safety issues, such as recommending inappropriate exercises that cause injury. Assessing and mitigating these risks is paramount.

Sprint 4: Social Accountability & Sharing

Duration: 2 weeks

Goals: Implement granular privacy-controlled sharing features, Develop a system for social support and accountability

Tasks: Build friend/connection system, Create shareable progress dashboards with customizable data visibility, Develop notification system for accountability buddies, Implement community features like comment threads on shared progress

-  Privacy and Data (#7 - Data)
-  Access to Data (#9 - Data)
-  Communication (#3 - Transparency)
-  Human Agency (#10 - Agency & Oversight)
-  Auditability (#18 - Accountability)

Sprint 5: Advanced Customization & Integration

Duration: 3 weeks

Goals: Deliver advanced features for expert users, Enable deeper integration with external tools and data

Tasks: Build advanced plan customizer (e.g., specific training splits, macro targets), Develop API integration with nutrition apps like MyFitnessPal, Enhance adaptation algorithm with more complex AI patterns, Implement data export functionality

- 📌 Explainability (#2 - Transparency)
- 📌 Privacy and Data (#7 - Data)
- 📌 Data Quality (#8 - Data)
- 📌 Access to Data (#9 - Data)
- 📌 Human Agency (#10 - Agency & Oversight)

Are the selected cards aligned with your project goals?

🚀 **Sprint 4: Social Accountability & Sharing**

Duration: 2 weeks

Goals: Implement granular privacy-controlled sharing features, Develop a system for social support and accountability

Tasks: Build friend/connection system, Create shareable progress dashboards with customizable data visibility, Develop notification system for accountability buddies, Implement community features like comment threads on shared progress

📌 Privacy and Data (#7 - Data)

Motivation: Privacy is a rising trend in the wake of various recent data misuse reveals. People are now increasingly conscious about handing out personal data. Similarly, regulations such as GDPR now affect data collection.

What to Do:

- Ask yourself:
- What data are used by the system?
- Does the system use or collect personal data? Why? How is the personal data used?
- Do you clearly inform your (end-)users about any personal data collection? E.g., ask for consent, provide an opportunity to revoke it etc.
- Have you taken measures to enhance (end-user) privacy, such as encryption or anonymization?
- Who makes the decisions regarding data use and collection? Do you have organizational policies for it?

Practical Example: Rather than collecting and selling data, appealing to privacy can also be profitable. Regulations are making it increasingly difficult to collect lots of personal data for profit. Privacy can be an alternate selling point in today's climate.

Justification: This card applies because the sprint involves sharing user biometric and progress data, raising significant privacy concerns. Users must be informed about data collection, consent must be obtained, and measures like encryption should be implemented to protect sensitive information.

📌 Access to Data (#9 - Data)

Motivation: Aside from carefully planning what data you collect and how it is also important to plan how it can or will be used and by whom.

What to Do:

- Ask yourself:
- Who can access the users' data, and under what circumstances?
- How do you ensure that the people who access the data: 1) have a valid reason to do so, and 2) adhere to the regulations and policies related to data?
- Do you keep logs of who accesses the data and when? Do the logs also tell why?
- Do you use existing data governance frameworks or protocols? Does your organization have its own?
- Who handles the data collection, storage and use?

Practical Example: Third parties you give access to the data can misuse it. A prominent example of this is the case of Cambridge Analytica and Facebook, in which data from Facebook was used questionably. However, such incidents can also paint your organization in a bad way.

Justification: This card is highly relevant as the sprint focuses on implementing granular privacy controls for data sharing. It necessitates defining who can access data (e.g., friends, family), under what conditions, and ensuring compliance with policies, including logging access for accountability.

Communication (#3 - Transparency)

Motivation: In practice, communication is a big part of being transparent with your stakeholders. Being transparent in communication can generate trust.

What to Do:

- Ask yourself:
- What is the goal of the system? Why is this particular system deployed in this specific area?
- What do you communicate about the system to its users and end-users? Is it enough for them to understand how the system works?
- If relevant to your system, do you somehow tell your (end-)users that they are interacting with an AI system and not with another human being?
- Do you collect user feedback? How is it used to change/improve the system?
- Are communication and transparency towards other audiences, such as the general public, relevant?

Practical Example: Clearly stating what data you collect and why can make you seem much more trustworthy. Compare this to a cellphone application that just states it needs to access your camera and storage.

Justification: This card applies because clear communication is essential for users to understand how their data is shared in the social features. It builds trust by informing users about data usage, obtaining consent, and potentially collecting feedback on sharing preferences.

Human Agency (#10 - Agency & Oversight)

Motivation: People interacting with the system or using it should be able to understand it sufficiently. Users should be able to make informed decisions based on its suggestions, or to challenge its suggestions. AI systems should let humans make independent choices.

What to Do:

- Ask yourself:
- Does the system interact with decisions by human actors, i.e. end users (e.g. recommending users actions or decisions, or presenting options)?
- Does the system communicate to its (end) users that a decision, content or outcome is the result of an algorithmic decision? Into how much detail does it go?
- In the system's use context, what tasks are done by the system and what tasks are done by humans?
- Have you taken measures to prevent overconfidence or overreliance on the system?

Practical Example: A medical system recommends diagnoses. How does the system communicate to doctors why it made a recommendation? How should the doctors know when to challenge the system? Does the system somehow change how patients and doctors interact?

Justification: This card is relevant because users must retain control over their data sharing decisions. The sprint involves features like customizable visibility and accountability buddies, requiring users to make informed choices without overrelying on automated suggestions.

Auditability (#18 - Accountability)

Motivation: Regulations affecting AI and data may necessitate audits of systems in the future. Similarly, if the system causes damage, an audit might be requested. It is good to have mechanisms in place beforehand.

What to Do:

- Ask yourself:

- Is the system auditable?
- Can an audit be conducted independently?
- Is the system available for inspection?
- What mechanisms facilitate the system's auditability? How is traceability and logging of the systems's processes and outcomes ensured?

Practical Example: In heavily regulated fields such as medicine, audits are typically required before a system can be utilized in the first place.

Justification: This card applies as the sprint includes implementing sharing features that require logging data access and changes. Auditability ensures that who accessed what data and when can be tracked, which is crucial for accountability and compliance in social sharing contexts.

Sprint 5: Advanced Customization & Integration

Duration: 3 weeks

Goals: Deliver advanced features for expert users, Enable deeper integration with external tools and data

Tasks: Build advanced plan customizer (e.g., specific training splits, macro targets), Develop API integration with nutrition apps like MyFitnessPal, Enhance adaptation algorithm with more complex AI patterns, Implement data export functionality

-  Explainability (#2 - Transparency)
-  Privacy and Data (#7 - Data)
-  Data Quality (#8 - Data)
-  Access to Data (#9 - Data)
-  Human Agency (#10 - Agency & Oversight)

Are the selected cards aligned with your project goals?

Sprint 5: Advanced Customization & Integration

Duration: 3 weeks

Goals: Deliver advanced features for expert users, Enable deeper integration with external tools and data

Tasks: Build advanced plan customizer (e.g., specific training splits, macro targets), Develop API integration with nutrition apps like MyFitnessPal, Enhance adaptation algorithm with more complex AI patterns, Implement data export functionality

-  Explainability (#2 - Transparency)

Motivation: If we cannot understand the reasons behind the actions of the AI, it is difficult to trust it.

What to Do:

- Ask yourself:
- Is explainability a goal for your system? How do you plan to ensure it?
- How well can each decision of the system be understood? By both developers and (end-)users?
- Did you try to use the simplest and most interpretable model possible for the context?
- Did you make trade-offs between explainability and accuracy? What kind of? Why?
- How familiar are you with your training or testing data? Can you change it when needed?
- If you utilize third party components in the system, how well do you understand them?

Practical Example: When interacting with a robot, users could ideally ask the robot 'why did you do that?' and receive an understandable response. This would make it much easier for them to trust a system.

Justification: The sprint involves enhancing the AI adaptation algorithm with more complex patterns, which could make the system's recommendations less transparent. Users and experts need to understand why specific workout/nutrition plans are generated, especially when customization options increase.

Privacy and Data (#7 - Data)

Motivation: Privacy is a rising trend in the wake of various recent data misuse reveals. People are now increasingly conscious about handing out personal data. Similarly, regulations such as GDPR now affect data collection.

What to Do:

- Ask yourself:
- What data are used by the system?
- Does the system use or collect personal data? Why? How is the personal data used?
- Do you clearly inform your (end-)users about any personal data collection? E.g., ask for consent, provide an opportunity to revoke it etc.
- Have you taken measures to enhance (end-user) privacy, such as encryption or anonymization?
- Who makes the decisions regarding data use and collection? Do you have organizational policies for it?

Practical Example: Rather than collecting and selling data, appealing to privacy can also be profitable. Regulations are making it increasingly difficult to collect lots of personal data for profit. Privacy can be an alternate selling point in today's climate.

Justification: The sprint includes API integration with external nutrition apps and data export functionality, which involves sharing and processing sensitive biometric and health data. This raises significant privacy concerns regarding data handling, consent, and compliance with regulations like GDPR.

Data Quality (#8 - Data)

Motivation: As AI are trained using data, the data used directly affects how the system operates. Both the nature and the quality and integrity of the data used has to align with goals of the system.

What to Do:

- Ask yourself:
- What are good or poor quality data in the context of your system?
- How do you evaluate the quality and integrity of your own data? Are there alternative ways?
- If you utilize data from external sources, how do you control their quality?
- Did you align your system with relevant standards (for example ISO, IEEE) or widely adopted protocols for daily data management and governance?
- How can you tell if your data sets have been hacked or otherwise compromised?

Practical Example: In 2017, Amazon scrapped its recruitment AI because of bad data. They used past recruitment data to teach the AI. As they had mostly hired men, the AI began to consider women undesirable based on the data.

Justification: Integrating with external apps like MyFitnessPal introduces third-party data sources whose quality and biases must be assessed. Poor data quality could lead to inaccurate or harmful fitness/nutrition recommendations, especially as the AI becomes more complex.

Access to Data (#9 - Data)

Motivation: Aside from carefully planning what data you collect and how it is also important to plan how it can or will be used and by whom.

What to Do:

- Ask yourself:
- Who can access the users' data, and under what circumstances?
- How do you ensure that the people who access the data: 1) have a valid reason to do so, and 2) adhere to the regulations and policies related to data?
- Do you keep logs of who accesses the data and when? Do the logs also tell why?

- Do you use existing data governance frameworks or protocols? Does your organization have its own?
- Who handles the data collection, storage and use?

Practical Example: Third parties you give access to the data can misuse it. A prominent example of this is the case of Cambridge Analytica and Facebook, in which data from Facebook was used questionably. However, such incidents can also paint your organization in a bad way.

Justification: The data export functionality and API integrations increase the risk of unauthorized data access. It's crucial to define and control who can access user data, especially when sharing with friends/family or external services, to prevent misuse.

Human Agency (#10 - Agency & Oversight)

Motivation: People interacting with the system or using it should be able to understand it sufficiently. Users should be able to make informed decisions based on its suggestions, or to challenge its suggestions. AI systems should let humans make independent choices.

What to Do:

- Ask yourself:
- Does the system interact with decisions by human actors, i.e. end users (e.g. recommending users actions or decisions, or presenting options)?
- Does the system communicate to its (end) users that a decision, content or outcome is the result of an algorithmic decision? Into how much detail does it go?
- In the system's use context, what tasks are done by the system and what tasks are done by humans?
- Have you taken measures to prevent overconfidence or overreliance on the system?

Practical Example: A medical system recommends diagnoses. How does the system communicate to doctors why it made a recommendation? How should the doctors know when to challenge the system? Does the system somehow change how patients and doctors interact?

Justification: The advanced customization features and AI enhancements could lead users to over-rely on automated recommendations. It's important to ensure users retain control and can modify or question the system's plans, particularly in a health-related context where personal needs vary.



Sprint 1: User Foundation & Exercise Library

Duration: 3 weeks

Goals: Create a comprehensive user onboarding process, Build a secure user profile database, Develop a searchable exercise library with video demonstrations and constraint tags

Tasks: Design and implement user onboarding questionnaire, Develop database schema for user profiles (biometrics, goals, injuries, restrictions), Create exercise database with video content, Tag exercises with modifiers (e.g., low-impact, knee-friendly, shoulder-safe)

Persona Discussion

Chloe Williams: As a user who wants to share my progress with friends for accountability, I'm concerned about how we communicate data collection. From ECCOLA Card 3 on Communication, are we clearly informing users during onboarding about what data is collected, why, and how it's shared? For example, if I share my weight loss journey, could someone misuse that information? We need to ensure transparency to build trust.

David Miller: Chloe, that's a valid point, but as a software engineer, I'm more worried about the backend. ECCOLA Card 7 on Privacy and Data asks about encryption and consent mechanisms. Are we implementing strong encryption for the user profile database? Biometric data is highly sensitive—imagine if it's leaked due to poor security. Also, how do we handle consent revocation? Users might change their minds about sharing data.

Robert Johnson: David, you're right about security, but let's not overlook accessibility. From Card 14, is the exercise library inclusive for someone like me with age-related limitations? The constraint tags like 'knee-friendly' need to be accurate; otherwise, I could get injured. Were older adults involved in testing this? Poor data quality here could lead to real harm, as Card 8 warns.

Sophia Garcia: Robert, I appreciate your concern—as a fitness instructor, I know bad exercise data can cause injuries. Card 8 on Data Quality is crucial: how are we ensuring the exercise videos and tags are high-quality? If we source from external providers, what standards do we use? For instance, an incorrect 'shoulder-safe' tag could worsen Ben's injury. We need rigorous validation processes.

Ben Carter: Sophia, exactly! And tying back to security from Card 12, if the exercise database is compromised, hackers could alter tags to cause harm. Have we assessed unique vulnerabilities, like data pollution attacks? As a physical therapist, I rely on accurate data for my own routine. Also, is the team including cybersecurity experts to prevent breaches?

Chloe Williams: Ben, that's a scary thought—data tampering could mislead users. But going back to communication, Card 3 asks if we collect feedback. How will users report issues with exercises or data inaccuracies? For accountability, we need a

simple way for people like me to provide input and see changes made. Otherwise, trust erodes quickly.

David Miller: Chloe, feedback mechanisms are key, but let's challenge that: what if feedback is biased? For example, if only advanced users like Sophia provide input, it might not represent beginners like Robert. Card 14 on Accessibility suggests involving diverse groups in development. Is our team representative? We might unintentionally exclude people with disabilities or different backgrounds.

Robert Johnson: David, you hit on a critical point. As an older adult, I might not be tech-savvy. Is the onboarding questionnaire easy to understand and accessible? Card 3 emphasizes communication—are we using plain language, not jargon? Also, for data privacy, Card 7: do we explain why we need biometric data in a way that's clear to everyone, not just engineers?

Sophia Garcia: Robert, that's important. But let's push further on data quality—Card 8 asks how we control quality from external sources. If we use AI to recommend exercises, could biased data lead to unfair recommendations? For instance, if most data is from young men, it might not suit Chloe or Robert. We need audits and diversity in data sourcing.

Ben Carter: Sophia, absolutely. And to tie it all together, security measures from Card 12 must include monitoring for data compromises. How do we detect if the exercise library has been hacked? Also, communication should include alerts for users if issues arise. Ultimately, we need a holistic approach: privacy, quality, security, and transparency working together to ensure ethical integrity.



Sprint 2: Core Plan Generation

Duration: 2 weeks

Goals: Develop the initial plan generation algorithm, Deliver basic workout and nutrition plans to the user

Tasks: Build algorithm to generate workout plans based on user profile, Build algorithm to generate basic nutrition guidelines based on profile, Display generated plans in the user interface, Implement basic print/save functionality

🗣️ Persona Discussion

David Miller: As a software engineer, I'm kicking off with Explainability from Card 2. How will we ensure users like me understand why the algorithm suggests specific workouts? For instance, if it recommends heavy weights for muscle gain, but I have

a busy schedule, I need to know the reasoning behind it to trust it. Can we make the algorithm's decisions transparent in the UI?

Chloe Williams: David, that's important, but let's not overlook Privacy and Data from Card 7. I share my progress with family, so how is my biometric data handled? Are we getting explicit consent before sharing anything, and is it encrypted? If not, I could be exposing sensitive health info without realizing it.

Robert Johnson: Chloe, you're right about privacy, but I'm more concerned with Human Agency from Card 10. At my age, I need to be able to challenge or modify plans if they feel unsafe. David, how explainable is it really? If the algorithm suggests something risky for my heart, can I easily see why and adjust it, or am I forced to follow it blindly?

Sophia Garcia: Robert, that ties into Data Quality from Card 8. As a fitness pro, I worry about the data input. If users like Ben enter incorrect injury data, the plans could be harmful. How do we validate data quality? David, is the algorithm trained on diverse data sets to avoid biases, like only recommending intense workouts for young people?

Ben Carter: Sophia, you hit on a key point. From Card 8, poor data quality could ruin my recovery. But back to Human Agency—David, how does the system prevent over-reliance? If I blindly follow a plan and aggravate my shoulder, who's responsible? We need clear disclaimers and user control features.

David Miller: Ben, good challenge. For Explainability, we could add tooltips explaining why exercises are chosen, like 'based on your age and goal.' But Chloe, on privacy, are we considering GDPR? We must encrypt data and get opt-in consent for sharing. Robert, for agency, we can include an 'edit plan' button prominently.

Chloe Williams: David, that sounds good, but what about Stakeholder Analysis from Card 0? Beyond users, how does this affect healthcare providers if plans go wrong? Or family members who see shared data? We need to map all stakeholders and their influences to avoid unintended harm.

Robert Johnson: Chloe, expanding on that, if my doctor sees a shared plan and disagrees, how does the system handle that? From Card 10, we need a way to incorporate external feedback. David, can the algorithm be updated based on professional input, or is it static?

Sophia Garcia: Robert, that's a great point. For Data Quality, if we allow professional input, how do we ensure it doesn't corrupt the algorithm with biased advice? David, are we using interpretable models so we can audit and refine them easily?

Ben Carter: Sophia, exactly. And back to Privacy—if data is shared with pros, is it anonymized? Chloe, have we considered scenarios where shared data leads to

social pressure or stigma? We need to test edge cases where privacy breaches affect mental health.

David Miller: Everyone, I think we've covered a lot. To refine: For Explainability, we'll add rationale displays. Privacy: implement encryption and granular consent. Data Quality: add data validation checks. Human Agency: enable plan customization and disclaimers. Stakeholders: document all, including indirect ones like insurers. Any final challenges?

Chloe Williams: David, one last thing—how do we ensure these features are accessible to non-tech users like me or Robert? If the UI is too complex, it defeats the purpose. We need simplicity in explanations and controls.

Robert Johnson: Agreed, Chloe. And let's not forget testing with diverse age groups to catch issues early. I think we've pushed this as far as we can ethically for now.



Sprint 3: Feedback & Adaptation Loop

Duration: 3 weeks

Goals: Enable users to log completed activities and provide feedback, Implement basic plan adaptation based on user input, Introduce reminder system for habit formation

Tasks: Develop workout and nutrition logging features, Create feedback mechanism (e.g., rate difficulty, log pain), Build logic for exercise substitution based on feedback, Implement simple notification system for reminders

💬 Persona Discussion

David Miller: As a software engineer, I'm concerned about the privacy aspects here. From ECCOLA Card 7, what biometric data are we collecting, and why? For instance, if we're logging heart rate or weight, how do we ensure users like me, who are busy and might not read fine print, are properly informed and consenting? Are we encrypting this data, or could it be misused?

Chloe Williams: David, that's a valid point, but as a marketing manager, I think communication is key here. Card 3 asks what we tell users about how the system works. If we're not clear about why we need their data and how feedback affects their plans, users like me—a busy mom—might not trust it. How do we balance transparency without overwhelming users with details?

Robert Johnson: I agree with both, but from my perspective as someone new to exercise, explainability is crucial. Card 2 says we need to understand why the system makes changes. If the app substitutes an exercise based on my feedback, I need to know why, especially at my age where safety is a concern. David, as a

developer, can you ensure the logic is simple enough for non-techies like me to grasp?

Sophia Garcia: Robert, that's important, but let's not forget data quality from Card 8. As a fitness instructor, I've seen bad data lead to poor recommendations. If users log incorrect info or the system has biases, it could harm someone's progress or even cause injury. David, how are we validating the input data, especially for advanced users like me who might push limits?

Ben Carter: All good points, but human agency from Card 10 is vital for me with my shoulder injury. The system might suggest adaptations, but users should always have the final say. For example, if it recommends an exercise that could aggravate my injury, I need to easily override it. How do we build in controls to prevent overreliance on automation?

David Miller: Ben, you're right—we need user override features. But Chloe, on communication: we could use simple pop-ups for consent and explain data usage in plain language. However, Robert, making the AI explainable might require trade-offs with accuracy. Should we use simpler models even if they're less precise, to ensure safety?

Chloe Williams: David, that's a trade-off worth discussing. But from a transparency angle, if we simplify too much, users might not get the best plans. Sophia, how would you feel if the system explains changes in basic terms, but advanced users like you can dive deeper? And Ben, for agency, maybe we add a 'why this change?' button with clear reasons.

Robert Johnson: I like that idea, Chloe. But going back to privacy, David: are we anonymizing data when sharing progress with friends? Card 7 mentions encryption—what specific measures? I wouldn't want my health data exposed accidentally.

Sophia Garcia: Robert, good point on sharing—we need opt-in controls for that. But on data quality, David, how do we handle cases where user feedback is inconsistent or biased? For instance, if someone rates an exercise too hard because they're having a bad day, it could skew the adaptation. Should we have safeguards like averaging feedback over time?

Ben Carter: Sophia, that ties into human agency—users might need education on providing accurate feedback. But David, from a development standpoint, are we following any standards for data governance, like ISO, as Card 8 suggests? Without that, we risk ethical lapses.

David Miller: Ben, we should adopt standards like ISO 27001 for data security. Chloe, for communication, we can implement tooltips or a help section that explains the AI's role. Robert, for explainability, we might use decision trees that are easier to

interpret, even if less accurate. But is that acceptable, or do we risk ineffective plans?

Chloe Williams: David, I think users would prefer slightly less accuracy if it means understanding and safety. Sophia, as an advanced user, would you agree? And for privacy, we must have clear consent forms and allow data deletion, as per GDPR. Robert, does that address your concerns?

Robert Johnson: Yes, but I'd want reminders to review privacy settings periodically. Also, on explainability, if the system changes my plan, it should say something like 'Based on your feedback that exercise X was too hard, we substituted Y which is safer for beginners.' Simple language helps.

Sophia Garcia: Robert, that works for basics, but for advanced users, we need more detail. Perhaps tiered explanations. Back to data quality: David, are we auditing the data for biases, like assuming all users have similar capabilities? That could exclude people with injuries like Ben's.

Ben Carter: Exactly, Sophia. The system must not make assumptions. David, how do we ensure the adaptation logic considers diverse user contexts? And for human agency, let's make sure users can easily revert changes or provide alternative feedback.

David Miller: All great points. We'll implement encryption, clear consent, tiered explanations, and user override features. Also, data quality checks for biases. I think we've covered the ethical angles from the cards. Any final challenges or refinements?

Chloe Williams: One last thing: communication about how feedback is used to improve the system over time. Users should know their input helps others, but anonymously. That builds trust and community, which is key for motivation.

Robert Johnson: Agreed. And ensure the notification system for reminders isn't intrusive—respect users' time and privacy.

Sophia Garcia: I think we've addressed the main issues. Let's prioritize these in development to avoid ethical pitfalls.

Ben Carter: Yes, with these measures in place, the system should be ethically sound for this sprint. No further improvements come to mind.



Sprint 4: Social Accountability & Sharing

Duration: 2 weeks

Goals: Implement granular privacy-controlled sharing features, Develop a system for social support and accountability

Tasks: Build friend/connection system, Create shareable progress dashboards with customizable data visibility, Develop notification system for accountability buddies, Implement community features like comment threads on shared progress

Persona Discussion

David Miller: Starting with ECCOLA Card 7 on Privacy and Data: As a software engineer, I'm concerned about what biometric data we're collecting and why. For instance, if we're gathering heart rate or weight, how is that used beyond generating plans? And are we encrypting this data to prevent breaches? Chloe, as a user who shares with family, do you think we're being transparent enough about data usage?

Chloe Williams: Good point, David. From Card 3 on Communication, I need clear, simple explanations about what data is shared and with whom. As a busy mom, I might not read fine print, so how do we ensure users like me understand and consent? Also, Robert, you're new to this—would you feel comfortable if the app didn't clearly explain how your data is handled?

Robert Johnson: Absolutely not, Chloe. Card 3 hits home for me—I need to know why the system is suggesting certain exercises and how my data is used. If I'm sharing progress with friends, I want granular control over what they see, as per Card 9 on Access to Data. Sophia, you share a lot; how do we prevent oversharing or misuse?

Sophia Garcia: Robert, that's crucial. From Card 9, we need logs of who accesses data and why. As a fitness pro, I'd want to share my journey but only with select people. David, how do we implement that technically without making it too complex? And Ben, with your injury background, shouldn't we have human oversight to challenge automated suggestions, as in Card 10 on Human Agency?

Ben Carter: Exactly, Sophia. Card 10 is key—the system might recommend workouts, but users like me with injuries need to override them easily. We must prevent overreliance. Also, from Card 19 on Ability to Redress, what if someone's data is misused or they face harassment through the sharing features? How do we handle complaints and offer compensation?

David Miller: Ben, that's a valid concern. For redress, we could implement a reporting system and clear policies. But Chloe, do you think users will actually use it, or will they just abandon the app if issues arise? We need to balance usability with safety.

Chloe Williams: David, from a marketing perspective, trust is everything. If users don't feel safe, they'll leave. So, for Card 7, we should anonymize data where possible and get explicit consent. Robert, how would you feel if the app asked for your permission each time you share something new?

Robert Johnson: I'd appreciate that, Chloe. It aligns with Card 3—clear communication. But Sophia, you might find it annoying if you share frequently. How can we make it efficient without sacrificing transparency?

Sophia Garcia: Robert, good point. Perhaps default settings with easy customization. From Card 9, we need to ensure that only authorized people access data, and logs are kept. Ben, do you think we should have a feature where accountability buddies are vetted or trained to avoid misuse?

Ben Carter: Sophia, that's an excellent idea. Incorporating human elements into the system, per Card 10, to maintain agency. Also, for Card 19, we must have a straightforward way for users to seek help if things go wrong, like a dedicated support channel. David, can we technically implement that without overloading the system?

David Miller: Yes, Ben, we can use existing frameworks for data governance and redress. But let's challenge this: what if the system's algorithms make a mistake in a workout plan, and someone gets hurt? How does Card 19 apply there? We need to define organizational responsibility clearly.

Chloe Williams: That's a scary thought, David. From Card 3, we must communicate the limitations of the system—it's not a substitute for professional advice. Robert, as someone with health concerns, would you want a disclaimer upfront?

Robert Johnson: Definitely, Chloe. And from Card 10, I'd want the ability to easily flag issues or suggest changes. Sophia, how do we ensure the community features don't become echo chambers or spread misinformation?

Sophia Garcia: Robert, that's where moderation and clear guidelines come in, tied to Card 9 on access control. We could have verified experts or filters. Ben, do you think we need a way for users to challenge algorithmic decisions publicly in the community?

Ben Carter: Sophia, that could empower users but also lead to conflict. From Card 10, we should provide tools for informed decision-making, not public debates that might cause harm. Perhaps private feedback loops instead. David, how do we log and review these interactions for improvement?

David Miller: Ben, we can implement audit logs as per Card 9. But let's refine: after this discussion, I think we've covered the need for transparency, control, and redress. Is there any aspect we're missing? For instance, what about data deletion rights under GDPR, from Card 7?

Chloe Williams: David, that's essential. Users should be able to delete their data and revoke consent easily. Robert, would you feel more trusting if the app offered that option prominently?

Robert Johnson: Yes, Chloe. And from Card 3, it should be communicated in plain language. I think we've addressed the key ethical points: privacy, access, communication, agency, and redress. Any final challenges?

Sophia Garcia: I agree, Robert. To sum up, we need granular privacy controls, clear communication, human oversight, and a robust redress system. Ben, do you see any gaps from a safety perspective?

Ben Carter: No, I think we've thoroughly covered it. The system should enhance human support without replacing it, and users must have avenues for recourse. This discussion has refined our approach well.



Sprint 5: Advanced Customization & Integration

Duration: 3 weeks

Goals: Deliver advanced features for expert users, Enable deeper integration with external tools and data

Tasks: Build advanced plan customizer (e.g., specific training splits, macro targets), Develop API integration with nutrition apps like MyFitnessPal, Enhance adaptation algorithm with more complex AI patterns, Implement data export functionality



Persona Discussion

David Miller: As a software engineer, I'm digging into Card 7 on Privacy and Data. With API integrations like MyFitnessPal, we're handling sensitive biometric data. How do we ensure we're not over-collecting? For example, are we only pulling necessary macros, or are we grabbing everything? And do we clearly inform users and get explicit consent before sharing data with third parties? Chloe, as someone who shares progress, wouldn't you want granular control over what's shared?

Chloe Williams: Absolutely, David. Card 9 on Access to Data is key for me too. When I share my progress with family, I need to know who else can access it—like developers or advertisers. Do we have logs for data access, and can users review them? But David, your point on consent: is a one-time consent enough, or should users be able to revoke it easily? Robert, since you're new to this, how would you feel about constant consent pop-ups?

Robert Johnson: Honestly, too many pop-ups might confuse me. But Card 2 on Explainability is more pressing. If the AI suggests a high-intensity workout, I need to know why, especially with my age. Can we have simple explanations, like 'This exercise improves heart health based on your doctor's advice'? Sophia, you're advanced—do you think detailed explanations would clutter the interface?

Sophia Garcia: Good point, Robert. For competitions, I need complex plans, but explanations must be clear. However, Card 8 on Data Quality worries me more. If we

integrate external data, how do we verify its accuracy? For instance, if MyFitnessPal has incorrect entries, our AI could suggest wrong macros. David, as an engineer, how would you implement data validation without slowing down the app?

Ben Carter: Jumping in with Card 10 on Human Agency. With my shoulder injury, I must be able to challenge AI recommendations. Does the system allow me to say 'This exercise hurts' and adjust the plan? But Sophia, your focus on data quality—if the data is biased towards young athletes, could it ignore older users like Robert? We need diversity in training data.

David Miller: Ben, that's a solid point on bias. For data quality, we could use checksums or API validation, but it might increase latency. Chloe, would you tolerate slower app performance for better data accuracy? And Robert, on explainability, perhaps we tier it—basic for beginners, detailed for experts. But does that risk oversimplifying for some?

Chloe Williams: I'd trade some speed for accuracy, especially if it affects my health. But back to privacy—Card 7 asks about encryption. Are we encrypting data at rest and in transit? And David, you mentioned consent revocation; how quickly can data be deleted if a user opts out? Robert, if you revoke consent, should the app still function with limited features?

Robert Johnson: If I revoke consent, I'd expect the app to still work but maybe not personalize plans. But Sophia, on data quality, how do we prevent scenarios like the Amazon AI bias? Should we audit external data sources regularly? Ben, with your therapy background, how can we ensure the AI doesn't recommend harmful exercises based on poor data?

Sophia Garcia: Regular audits are a must, and we need transparency in data sources. But Ben, your human agency point—can we add a 'flag for review' feature where users report issues, and humans step in? However, that might not scale. David, is there a way to automate bias detection in the AI without compromising performance?

Ben Carter: Automation is fine, but Card 10 emphasizes human oversight. Perhaps a hybrid model where critical decisions are reviewed by professionals. But Chloe, on access logs—if we log everything, could that itself become a privacy risk? How do we balance transparency with data minimization?

David Miller: That's a tough one. We could anonymize logs or use aggregated data for analysis. But after this discussion, I think we've covered the ethical angles: need for clear consent, data quality checks, explainable AI, controlled access, and user agency. Any counterarguments or can we refine no further? For instance, is there a scenario where full automation is unethical regardless of safeguards?

Chloe Williams: In cases of medical conditions, full automation might always be risky without human input. But for general fitness, perhaps it's acceptable with opt-outs. I think we've pushed this as far as we can—the key is implementing these ideas with user-centric design.

Robert Johnson: Agreed. As long as I can understand and control my plan, I'm satisfied. No further improvements from me.

Sophia Garcia: Same here. The ethical framework seems solid if executed well.

Ben Carter: I'll echo that. We've addressed the major concerns, and further refinement would be implementation-specific.

Ethical Requirements

Discussion

David Miller: As we refine, I'm revisiting Card 7: we must implement data minimization—only collect biometrics essential for personalized plans, like age, weight, and goals, not extraneous data. Chloe, does this align with your need for transparency? And for encryption, we'll use AES-256 for data at rest and TLS 1.3 in transit. But Robert, if we minimize data, could that limit plan accuracy for edge cases?

Chloe Williams: David, minimization is good, but transparency means explaining why each data point is needed—e.g., 'We use your weight to calculate calorie needs.' For sharing, granular controls are key: let users choose what to share (e.g., progress photos but not weight). Sophia, as an influencer, would you want options to share selectively with different groups, like public vs. close friends?

Robert Johnson: Chloe, that's smart. But safety first: from Card 14, we need age-specific exercise modifiers—e.g., low-impact alternatives automatically offered for users over 60. David, can the algorithm tag exercises with safety levels? And Ben, should there be mandatory disclaimers for users with health conditions, like 'Consult your doctor before starting'?

Sophia Garcia: Robert, yes—and for data quality (Card 8), we must audit external sources like MyFitnessPal for accuracy. David, can we implement real-time validation, like flagging inconsistent entries? But Chloe, if we do, will users find it intrusive? Also, for advanced users, I want the option to input custom exercises with verified form videos to prevent misinformation.

Ben Carter: Sophia, custom inputs need moderation to avoid harmful content. From Card 10, users must easily report issues—e.g., a 'flag unsafe exercise' button with quick human review. David, technically, can we prioritize these reports? And Robert,

for redress (Card 19), we need a clear process: if someone gets injured, offer support like connecting to a health professional, not just a refund.

David Miller: Ben, we can implement a triage system for reports. But Chloe, will users trust it? Perhaps we display moderation status publicly, like 'Verified by trainers.' Sophia, does that address your concern about data quality? And for minimization, we'll anonymize data used for analytics, but is that enough for GDPR compliance?

Chloe Williams: David, public verification builds trust. But for privacy, users should opt in to data use for analytics—no pre-checked boxes. Robert, would you prefer periodic privacy check-ins, like a quarterly reminder to review settings? Also, for sharing, we need to prevent social pressure: options to hide metrics like weight if users feel judged.

Robert Johnson: Chloe, quarterly reminders are great. And for safety, the app should never force exercises—always allow skips or substitutions. Ben, should we integrate with wearables for real-time feedback, like heart rate alerts during workouts? But David, that means more data—how do we ensure it's not misused?

Sophia Garcia: Robert, wearable integration is useful but risky. We need explicit consent for each data stream. For data quality, let's add user ratings for exercises, weighted by experience level—e.g., a beginner's 'too hard' rating triggers fewer changes than a pro's. David, can we implement that without bias?

Ben Carter: Sophia, weighting ratings could work, but we must avoid elitism. From Card 14, ensure diverse testing groups include all ages and abilities. David, for redress, have a dedicated email support with a 24-hour response time for urgent issues. Finally, all features must be WCAG-compliant for accessibility—Robert, would large font options help?

David Miller: Ben, WCAG compliance is non-negotiable. I think we've covered it: minimize data, encrypt everything, granular sharing, safety tags, data validation, human moderation, and robust redress. Any final gaps? For instance, if the AI adapts plans, should we log all changes for auditability?

Chloe Williams: Yes, audit logs are essential for transparency. But keep them user-accessible. I think we're done—this balances innovation with ethics.

Ethical Requirements

- Implement data minimization: collect only essential biometric data (e.g., age, weight, goals) with clear explanations for each data point.
- Use AES-256 encryption for data at rest and TLS 1.3 for data in transit to protect user privacy.

- Provide granular sharing controls: users can choose what data to share (e.g., progress photos, metrics) and with whom (e.g., public, friends, family).
- Include age-specific safety modifiers for exercises (e.g., automatic low-impact alternatives for users over 60) and mandatory disclaimers for users with health conditions.
- Validate external data sources (e.g., MyFitnessPal) in real-time and allow user-reported inaccuracies with prioritization for safety issues.
- Enable custom exercise inputs with moderation and verification to prevent harmful or misinformation.
- Implement a 'flag unsafe content' feature with human review and a clear redress process (e.g., support escalation, professional connections).
- Require explicit opt-in consent for data use in analytics and wearable integrations, with no pre-checked boxes.
- Offer periodic privacy setting reminders (e.g., quarterly) and options to hide sensitive metrics to reduce social pressure.
- Ensure WCAG compliance for accessibility, including features like adjustable font sizes and voice commands.
- Maintain audit logs for AI-driven plan changes and data access, accessible to users upon request.
- Weight user feedback by experience level to improve data quality without introducing bias, and test features with diverse user groups.

Ethical User Stories

Personas Discussion

David Miller: As a software engineer, I'm big on data privacy. For the first requirement on data minimization, I think we should have a user story like: 'As a user, I want the app to collect only essential biometric data with clear opt-ins for more, so that my privacy is protected and I avoid unnecessary data exposure.' This aligns with my need for efficiency and control, especially with my hectic schedule.

Chloe Williams: That makes sense, David. As a busy mom, I appreciate simplicity. For the tiered explanations, how about: 'As a user, I want explanations for recommendations tailored to my knowledge level, so that I can understand my plan without getting overwhelmed.' This helps me stay motivated without confusion.

Robert Johnson: I agree, clarity is key for someone like me who's new to this. But let's not forget safety. For the mandatory disclaimers, I'd say: 'As a user, I want prompts to consult a doctor when risks are detected, so that I can exercise safely given my age and health conditions.' Safety first, always.

Sophia Garcia: Good points, everyone. As a fitness pro, I value customization and sharing. For granular data sharing: 'As a user, I want to choose exactly what data to share and with whom, so that I can inspire others without compromising my privacy.' This lets me control my social presence.

Ben Carter: Speaking of safety, as a physical therapist, I'd emphasize adaptive onboarding. How about: 'As a user, I want the app to ask for injury details only when relevant, with clear reasons, so that my plan is safe and accurate without prying.' This prevents unnecessary data collection and focuses on what's needed.

David Miller: Building on that, for the redress system, we need: 'As a user, I want an easy way to report issues and get human help, so that problems are fixed fairly and the app improves ethically.' This ensures accountability, which is crucial in software.

Chloe Williams: And for biases, we should have: 'As a user, I want the app to be audited for fairness and be transparent about data use, so that it works well for people of all ages and abilities.' This promotes inclusivity, which matters for my diverse social circle.

Robert Johnson: Lastly, on balancing free and premium features: 'As a user, I want core safety and accessibility features to be free, so that everyone can benefit without ethical compromises.' This makes it accessible for retirees like me on a budget.

Sophia Garcia: I think we've covered all the requirements. Let's refine these stories to ensure they're clear and actionable. Everyone's input has made them stronger.

Ben Carter: Agreed. This collaborative approach ensures the stories are well-rounded and address real user needs from different perspectives.

Generated Ethical User Stories

- As a user, I want the app to collect only essential biometric data with clear opt-ins for additional data, so that my privacy is respected and I have control over my information.
- As a user, I want adjustable explanations for AI-generated recommendations based on my expertise level, so that I can understand the advice in a way that suits me, promoting transparency and education.
- As a user, I want granular control over which metrics to share and with whom, so that I can maintain privacy while engaging socially with friends and family.
- As a user, I want adaptive onboarding that prompts for additional context like injuries only when needed, with clear communication on why it improves accuracy, so that my plan is tailored without unnecessary data collection.

- As a user, I want mandatory safety disclaimers and prompts to consult healthcare professionals when risks are detected, so that my health and safety are prioritized.
- As a user, I want an easy-to-use redress system with reporting features and options for human review, so that issues are addressed fairly and the app improves continuously without privacy breaches.
- As a user, I want the app to regularly audit data sources and algorithms for biases, with transparency in handling, so that it is inclusive and fair across all fitness levels and abilities.
- As a user, I want core safety and accessibility features to be available for free, so that ethical standards are maintained and the app is accessible to everyone without a two-tier system.